

1 Introduction

Knowledge editing changes what a language model believes about the world, and evaluating it requires benchmarks whose metrics are defined precisely enough that two papers reporting the same number mean the same thing. This paragraph is padding so that the two columns fill and a sentence has to cross from the first column into the second one, which is the case the scanner must survive without splitting the sentence in two. We repeat the point once more to push the text further down the first column of the first page of this short document.

Zhong et al. (2023) define multi-hop accuracy as correct if any of three paraphrased questions is answered. MQuAKE (Zhong et al., 2023) evaluates whether an edit propagates to its multi-hop consequences. Temporal benchmarks (Jang et al., 2022; Zhong et al., 2023) track how facts change over time. Mutual information follows the definition in Cover and Thomas (2006). The scaling result of Smith et al. (2020a) was later revisited under distribution shift (Smith et al., 2020b). We follow the evaluation protocol of Bao et al. (2023) for all ranking runs. A claim attributed to an unknown source (Nobody et al., 2019) has no entry in the list.

References

Keqin Bao, Jizhi Zhang, Yang Zhang, Wenjie Wang, Fuli Feng, and Xiangnan He. TALL-Rec: An effective and efficient tuning framework to align large language model with recommendation. In *Proceedings of RecSys*, 2023. doi: 10.1145/3604915.3608857.

Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley-Interscience, 2006.

Joel Jang, Seonghyeon Ye, Changho Lee, Sohee Yang, Joongbo Shin, Janghoon Han, Gyeonghun Kim, and Minjoon Seo. TemporalWiki: A lifelong benchmark for training and evaluating ever-evolving language models. In *Proceedings of EMNLP*, 2022.

Alice Smith, Bob Brown, and Carol Green. Scaling

laws for retrieval augmented generation. *Transactions on Machine Learning Research*, 2020a.

Alice Smith, Bob Brown, and Dan White. Calibrating retrieval augmented generation under distribution shift. *Journal of Machine Learning Research*, 2020b.

Zexuan Zhong, Zhengxuan Wu, Christopher D. Manning, Christopher Potts, and Danqi Chen. MQuAKE: Assessing knowledge editing in language models via multi-hop questions. In *Proceedings of EMNLP*, 2023. arXiv:2305.14795.

A Extra Results

The appendix repeats the protocol of Bao et al. (2023) on a second dataset.