

1 Introduction

Knowledge editing changes what a language model believes about the world, and evaluating it requires benchmarks whose metrics are defined precisely enough that two papers reporting the same number mean the same thing. This paragraph is padding so that the two columns fill and a sentence has to cross from the first column into the second one, which is the case the scanner must survive without splitting the sentence in two. We repeat the point once more to push the text further down the first column of the first page of this short document.

MQuAKE [1] evaluates whether an edit propagates to its multi-hop consequences. Earlier benchmarks [2–4] track how facts change over time. We follow the evaluation protocol of [5] for all ranking runs. Scores lie in the interval $[0, 1]$ and the median was 3.3 [2,5] across seeds. For the vector we set $c_0 = [1]$ before the first round.

References

- [1] Z. Zhong, Z. Wu, C. D. Manning, C. Potts, and D. Chen, “MQuAKE: Assessing knowledge editing in language models via multi-hop questions,” in *Proceedings of EMNLP*, 2023. arXiv:2305.14795.
- [2] J. Jang, S. Ye, C. Lee, S. Yang, J. Shin, J. Han, G. Kim, and M. Seo, “TemporalWiki: A life-long benchmark for training and evaluating ever-evolving language models,” in *Proceedings of EMNLP*, 2022.
- [3] T. M. Cover and J. A. Thomas, *Elements of Information Theory*. Wiley-Interscience, 2006.
- [4] A. Smith, B. Brown, and C. Green, “Scaling laws for retrieval augmented generation,” *Transactions on Machine Learning Research*, 2020.
- [5] K. Bao, J. Zhang, Y. Zhang, W. Wang, F. Feng, and X. He, “TALLRec: An effective and efficient tuning framework to align large language model with recommendation,” in *Proceedings of RecSys*, 2023.